

Explore the different object attributes of the dataframe

- Pandas is the high-powered Python library used for data analysis.
 - It is built on top of NumPy and is capable of handling large datasets efficiently.
 - Pandas is also used for data visualization as it can plot data in a tabular or graphical form.
- It will be the basis of GeoPandas, which we will use to look at vector data
 - For now, we will start with an ordinary table of data.

```
# Do our imports
import pandas
import os

# Set our file paths (YOU SHOULD HAVE DOWNLOADED THIS FROM GOOGLE DRIVE!)
data_directory = '../.../data'
food_prices_filename = 'world_monthly_food_prices.csv'
food_prices_path = os.path.join(data_directory, food_prices_filename)

# Pandas provides a read_csv function that will read a CSV file into a dataframe
food_prices = pandas.read_csv(food_prices_path)

# Print the dataframe using default printing options
print(food_prices)
```

	Domain	Code	Domain
0	CP	Consumer Price Indices	\
1	CP	Consumer Price Indices	
2	CP	Consumer Price Indices	
3	CP	Consumer Price Indices	
4	CP	Consumer Price Indices	

```

..
176          CP  Consumer Price Indices
177          CP  Consumer Price Indices
178          CP  Consumer Price Indices
179          CP  Consumer Price Indices
180  FAOSTAT Date: Wed Aug 17 16:12:56 CEST 2016          NaN

```

	AreaCode	AreaName	ElementCode	ElementName	ItemCode
0	5000.0	World	7001.0	January	23013.0 \
1	5000.0	World	7001.0	January	23013.0
2	5000.0	World	7001.0	January	23013.0
3	5000.0	World	7001.0	January	23013.0
4	5000.0	World	7001.0	January	23013.0
..
176	5000.0	World	7012.0	December	23013.0
177	5000.0	World	7012.0	December	23013.0
178	5000.0	World	7012.0	December	23013.0
179	5000.0	World	7012.0	December	23013.0
180	NaN	NaN	NaN	NaN	NaN

	ItemName	Year	Value	Flag
0	Consumer Prices, Food Indices (2000 = 100)	2000.0	99.5	NaN \
1	Consumer Prices, Food Indices (2000 = 100)	2001.0	101.5	NaN
2	Consumer Prices, Food Indices (2000 = 100)	2002.0	106.5	NaN
3	Consumer Prices, Food Indices (2000 = 100)	2003.0	112.5	NaN
4	Consumer Prices, Food Indices (2000 = 100)	2004.0	119.2	NaN
..
176	Consumer Prices, Food Indices (2000 = 100)	2011.0	210.3	NaN
177	Consumer Prices, Food Indices (2000 = 100)	2012.0	224.2	NaN
178	Consumer Prices, Food Indices (2000 = 100)	2013.0	239.0	NaN
179	Consumer Prices, Food Indices (2000 = 100)	2014.0	250.9	NaN
180	NaN	NaN	NaN	NaN

	FlagD
0	Official data
1	Official data
2	Official data
3	Official data
4	Official data
..	...
176	Official data
177	Official data
178	Official data

```
179 Official data
180           NaN
```

```
[181 rows x 12 columns]
```

- For instance look at the column names

```
print('List of column names:', food_prices.columns)
```

```
List of column names: Index(['Domain Code', 'Domain', 'AreaCode', 'AreaName', 'ElementCode',
                             'ElementName', 'ItemCode', 'ItemName', 'Year', 'Value', 'Flag',
                             'FlagD'],
                             dtype='object')
```

Get a specific column

- To get a column, use square braces with the name of the column.

```
food_prices_value_column = food_prices['Value']

print('Specific column:', food_prices_value_column)
```

```
Specific column: 0          99.5
1          101.5
2          106.5
3          112.5
4          119.2
...
176        210.3
177        224.2
178        239.0
179        250.9
180         NaN
Name: Value, Length: 181, dtype: float64
```

```
print('Specific value in that column:', food_prices_value_column[6])
```

```
Specific value in that column: 132.2
```

Under the hood

- All parts of the pandas dataframe are represented via numpy arrays
 - If you want to access them as numpy arrays instead of Pandas dataframes or columns, you can use the `.values` attribute


```
numpy_array = food_prices_value_column.values
print(numpy_array)
```

```
[ 99.5 101.5 106.5 112.5 119.2 125.   132.2 139.7 155.5 172.   184.2 199.9
 213.2 227.2 240.7 100.5 101.9 107.5 113.7 119.5 126.8 133.   141.2 159.4
 172.6 185.9 201.6 213.7 229.7 241.8   99.9 102.3 107.2 113.9 120.   126.9
 132.9 141.6 161.   173.   185.6 201.9 214.9 229.2 242.3   99.7 103.1 107.4
 114.2 120.9 127.3 133.5 141.9 163.   173.4 186.1 202.6 215.6 230.6 243.1
   99.7 103.5 107.7 114.2 121.4 127.4 134.2 142.4 164.3 174.   186.1 203.2
 215.9 230.6 243.8   99.3 103.2 107.6 113.7 121.6 127.   134.3 143.7 165.8
 174.3 186.4 204.2 216.   232.1 244.6   99.2 103.3 107.7 113.5 121.7 128.
 134.3 145.3 167.5 175.9 188.   205.7 217.3 234.   246.4   99.6 103.9 108.5
 113.9 122.8 128.4 135.2 147.4 168.6 177.5 189.9 207.1 219.5 236.2 247.7
 100.1 104.4 109.4 114.9 123.9 129.   136.7 149.   169.7 178.8 192.1 208.5
 221.   237.4 249.1 100.2 105.   109.6 116.   124.5 129.8 137.6 150.4 170.1
 179.3 193.5 209.4 221.4 238.4 249.6 100.6 105.3 110.6 117.2 124.5 130.3
 138.5 151.5 169.8 180.4 195.3 209.6 222.1 239.4 249.9 100.9 105.7 111.3
 118.   124.6 130.5 140.1 153.2 170.   182.2 197.3 210.3 224.2 239.   250.9
  nan]
```

Default plotting

- By default, pandas has a plot function
 - Let's try it

```
import matplotlib
from matplotlib import pyplot as plt
plt.plot(food_prices['Value'])
plt.show()
```



Discussion point

- Why is the data like this?

Pandas Functionality

- Let's start by importing everything afresh and setting a random seed.
 - Setting the random seed means we will get the same random numbers each time we run the code.

```
import os

import numpy as np
import pandas as pd

# Set a seed value for the random generator
np.random.seed(48151623)
```

Create an individual column

- In pandas, they are called Series
- One very important thing to note is that in the output, we see TWO columns
 - Pandas always includes an index for every Series or Dataframe
 - The index is the leftmost column, and is used to identify each row

```
# Creating a Series by passing a list of values, letting pandas create a default integer index
s = pd.Series([1, 3, 5, np.nan, 6, 8])
print(s)
```

```
0    1.0
1    3.0
2    5.0
3    NaN
4    6.0
5    8.0
dtype: float64
```

There are many built-in ways of generating series

- For instance, you can automatically generate a series of dates

```
# Pandas is very detailed in dealing with dates and all the quirks (leap year?) that this
dates = pd.date_range('20130101', periods=6)
print(dates)
```

```
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')
```

```
# Creating a DataFrame by passing a NumPy array, with a datetime index and labeled columns
df = pd.DataFrame(np.random.randn(6, 4), columns=list('ABCD'))
print('df:\n', df)
```

```
df2 = pd.DataFrame({'A': 1.,
                    'B': pd.Timestamp('20130102'),
                    'C': pd.Series(1, index=list(range(4)), dtype='float32'),
```

```

'D': np.array([3] * 4, dtype='int32'),
'E': pd.Categorical(["test", "train", "test", "train"]),
'F': 'foo'})

# df.head()
# print(df.index)
# print(df.columns)
df.describe()

# Also note that a dataframe is really just a numpy array dressed up with extra trappings.
# can get back the raw array (though this might lose a lot of functionality).
a = df.to_numpy()
print('a\n', a)

# Sorting Values:

# Also, I want to illustrate THE MOST COMMON MISTAKE people make with Pandas.

# The sort_values method (a method is just a function attached to an object) returns a NEW
# Thus, in the line below, if you just printed df, it would not be sorted because we didn't
df.sort_values(by='B')
# print('Not sorted:\n', df)

# Easy way to get around this is just to assign the returned dataframe to a variable (even if it's the same object)
df = df.sort_values(by='B')
# print('Sorted with return:\n', df)

# Alternatively, if you hate returning things, there is the inplace=True command, which will modify the original dataframe
df.sort_values(by='B', inplace=True)
# print('Sorted inplace:\n', df)

## Selection/subsetting of data

# Selecting a single column, which yields a Series, equivalent to df.A
df['A']
df.A

# Selecting via [], which slices the rows.
df[0:3] # CAN BE SLOW

# Note, slicing above, which uses the

```

```

# standard Python / Numpy expressions for selecting and setting are intuitive its best to use
# the optimized pandas data access methods, .at, .iat, .loc and .iloc.

## Selecting by LABELS, loc and iloc

r = df.loc[0] # 0-th row.

# print('r', r)

# Discuss difference between df['A'] and df.loc[0]
r = df.loc[0, 'A']

r = df.loc[:, 'A'] # Colon is a slice, an empty colon means ALL the values.

# OPTIMIZATION:
# for faster single point access, use:
r = df.at[0, 'A']

# SELECTING BY POSITION
r = df.iloc[3]

# Selecting with slices
r = df.iloc[3:5, 0:2]

# Slices again with an empty slice.
r = df.iloc[1:3, :]

r = df.iloc[:, 1:3]

# SIMILAR OPTIMIZATION:
r = df.iat[1, 1]

# Boolean indexing
# Using a single column's values to select data.
r = df[df['A'] > 0]

# Make a copy (why?) and add a column
df2 = df.copy()
df2['E'] = ['one', 'one', 'two', 'three', 'four', 'three']
r = df2[df2['E'].isin(['two', 'four'])]

```



```

# Setting by assigning with a NumPy array:
df.loc[:, 'D'] = np.array([5] * len(df))

# Missing data

# First we're going to create a new df by "reindexing" the old one, which will shuffle the
# order according to the index provided. At the same time, we're going to add on a new, em
# EE, which we set as 1 for the first two obs.

df1 = df.reindex(index=[2, 0, 1, 3], columns=list(df.columns) + ['EE'])
df1.loc[0:1, 'EE'] = 1
# print(df1)

# Apply: Similar to R. Applies a function across many cells (fast because it's vectorized)
df.apply(np.cumsum)
df.apply(lambda x: x.max() - x.min())

# Concat
s = pd.Series(range(0, 6))
# print('s', s)

r = pd.concat([df, s]) # Concatenate it, default is by row, which just puts it on the bott
r = pd.concat([df, s], axis=1) # Concatenate as a new column

# print(r) # Result when concatenating a series of the same size.

s = pd.Series(range(0, 7))
r = pd.concat([df, s], axis=1) # Concatenate as a new column

s = pd.Series(range(0, 2))
r = pd.concat([df, s], axis=1) # Concatenate as a new column

# Join
# SQL style merges. See the Database style joining section.

left = pd.DataFrame({'key': ['foo', 'bar'], 'lval': [1, 2]})
right = pd.DataFrame({'key': ['foo', 'bar'], 'rval': [4, 5]})

# print(left)
# print(right)

```

```

df = pd.merge(left, right, on='key')

# print('df:\n', df)

# Stacking
stacked = df.stack()
# print('stacked:\n', stacked)

# Pivot Tables
df = pd.DataFrame({'A': ['one', 'one', 'two', 'three'] * 3,
                   'B': ['A', 'B', 'C'] * 4,
                   'C': ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 2,
                   'D': np.random.randn(12),
                   'E': np.random.randn(12)})

# print(df) # SPREADSHEET VIEW
df = pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
# print(df) # Multiindexed (Pivot table) view.

# NOTICE that a pivot table is just the above data but where specific things have been made
# indices.

# PLOTTING
ts = pd.Series(np.random.randn(1000),
               index=pd.date_range('1/1/2000', periods=1000))

ts = ts.cumsum()
ts.plot()
import matplotlib.pyplot as plt
# plt.show()

# Writing to files

df.to_csv('foo.csv')

# Reading files:

# FIRST NOTE, here we are using relative paths (which you should almost always do too). This
# this path works if you organized your data into the folder structure I suggested.

```

```

wdi_filename = "WDI_CO2_data.csv"
wdi_path = os.path.join(data_directory, wdi_filename)
df = pd.read_csv(wdi_path)

print('csv read as a df\n', df)

# For reference, here's the Excel version
# df = pd.read_excel('foo.xlsx', 'Sheet1', index_col=None, na_values=['NA'])

cols = list(df.columns)

# Make a subset of only 2 cols
r = df[['Country Code', '1970 [YR1970]']]
# print(r)

r = df.loc[df['Country Code'] == 'CAN']
# print('r', r)

rr = r.loc[df['Series Name'] == 'Total greenhouse gas emissions (kt of CO2 equivalent)']
print(rr)

# Class exercise: Plot the emissions of CO2 for Canada (or wherever I don't care).

```

df:

	A	B	C	D
0	0.336581	0.923754	-0.277124	0.388604
1	-1.295428	3.296657	-0.698246	0.245552
2	-1.086536	1.187113	-0.153344	-1.264476
3	0.798694	1.330577	-0.113975	-0.060949
4	-0.076321	-0.240310	0.728421	-0.384309
5	0.634212	1.605129	1.415844	0.385849

a

```

[[ 0.33658127  0.92375415 -0.27712413  0.38860406]
 [-1.2954285   3.29665716 -0.69824576  0.24555238]
 [-1.0865361   1.18711302 -0.15334368 -1.26447611]
 [ 0.79869365  1.33057707 -0.11397528 -0.06094946]
 [-0.07632075 -0.24031028  0.72842085 -0.38430931]
 [ 0.6342119   1.60512881  1.41584369  0.38584939]]

```

csv read as a df

	Country Name	Country Code
0	Afghanistan	AFG \
1	Afghanistan	AFG

2	Afghanistan	AFG
3	Afghanistan	AFG
4	Afghanistan	AFG
...
2640	NaN	NaN
2641	NaN	NaN
2642	NaN	NaN
2643	Data from database: World Development Indicators	NaN
2644	Last Updated: 10/15/2020	NaN

	Series Name	Series Code
0	Access to clean fuels and technologies for coo...	EG.CFT.ACCS.ZS \
1	Access to electricity (% of population)	EG.ELC.ACCS.ZS
2	Adjusted net enrollment rate, primary (% of pr...	SE.PRM.TENR
3	Arable land (% of land area)	AG.LND.ARBL.ZS
4	Agricultural methane emissions (thousand metri...	EN.ATM.METH.AG.KT.CE
...
2640	NaN	NaN
2641	NaN	NaN
2642	NaN	NaN
2643	NaN	NaN
2644	NaN	NaN

	1960 [YR1960]	1961 [YR1961]	1962 [YR1962]	1963 [YR1963]
0 \
1
2
3	..	11.7176730079956	11.7942591060871	11.8708452041785
4
...
2640	NaN	NaN	NaN	NaN
2641	NaN	NaN	NaN	NaN
2642	NaN	NaN	NaN	NaN
2643	NaN	NaN	NaN	NaN
2644	NaN	NaN	NaN	NaN

	1964 [YR1964]	1965 [YR1965]	...	2011 [YR2011]	2012 [YR2012]
0	22.33	24.08 \
1	43.2220189082037	69.1
2
3	11.94743130227	11.94743130227	...	11.9336458046135	11.9321140826517
4
...

2640	NaN	NaN	...	NaN	NaN
2641	NaN	NaN	...	NaN	NaN
2642	NaN	NaN	...	NaN	NaN
2643	NaN	NaN	...	NaN	NaN
2644	NaN	NaN	...	NaN	NaN

	2013 [YR2013]	2014 [YR2014]	2015 [YR2015]	2016 [YR2016]	
0	26.17	27.99	30.1	32.44	\
1	68.9332656860352	89.5	71.5	97.7	
2	
3	11.9244554728426	11.903011365377	11.893821033606	11.8386790429801	
4	
...	
2640	NaN	NaN	NaN	NaN	
2641	NaN	NaN	NaN	NaN	
2642	NaN	NaN	NaN	NaN	
2643	NaN	NaN	NaN	NaN	
2644	NaN	NaN	NaN	NaN	

	2017 [YR2017]	2018 [YR2018]	2019 [YR2019]	2020 [YR2020]	
0	
1	97.7	98.7132034301758	
2	
3	
4	
...	
2640	NaN	NaN	NaN	NaN	
2641	NaN	NaN	NaN	NaN	
2642	NaN	NaN	NaN	NaN	
2643	NaN	NaN	NaN	NaN	
2644	NaN	NaN	NaN	NaN	

[2645 rows x 65 columns]

	Country Name	Country Code	
369	Canada	CAN	\

	Series Name	Series Code	
369	Total greenhouse gas emissions (kt of CO2 equi...	EN.ATM.GHGT.KT.CE	\

	1960 [YR1960]	1961 [YR1961]	1962 [YR1962]	1963 [YR1963]	1964 [YR1964]	
369	\

	1965 [YR1965]	...	2011 [YR2011]	2012 [YR2012]	2013 [YR2013]	
--	---------------	-----	---------------	---------------	---------------	--

```

369          .. ... 1033481.98200961 1027063.85487082          .. \
      2014 [YR2014] 2015 [YR2015] 2016 [YR2016] 2017 [YR2017] 2018 [YR2018]
369          ..          ..          ..          ..          .. \
      2019 [YR2019] 2020 [YR2020]
369          ..          ..

[1 rows x 65 columns]

```

